

S.NO	DATE	NAME OF THE EXPERIMENT	PAGE.NO	STAFF'S SIGNATURE
1.		Install Virtualbox/VMware/equivalent open source cloud Workstation with different flavours of Linux or Windows OS on top of windows 8 and above.	1	
2.		Install a C compiler in the virtual machine created using a virtual box and execute Simple Programs.	8	
3.		Install Google App Engine. Create a hello world app and other simple web applications using python/java.	10	
4.		Use the GAE launcher to launch the web applications.	13	
5.		Simulate a cloud scenario using CloudSim and run a scheduling algorithm that is not present in CloudSim.	16	
6.		Find a procedure to transfer the files from one virtual machine to another virtual machine.	21	
7.		Install Hadoop single node cluster and run simple applications like wordcount.	25	
8.		Creating and Executing Your First Container Using Docker.	36	
9.		Run a Container from Docker Hub.	38	

**EX.NO:1**

## **INSTALLATION OF VIRTUALBOX WITH DIFFERENT FLAVOURS OF LINUX/WINDOWS**

**DATE:**

**AIM:**

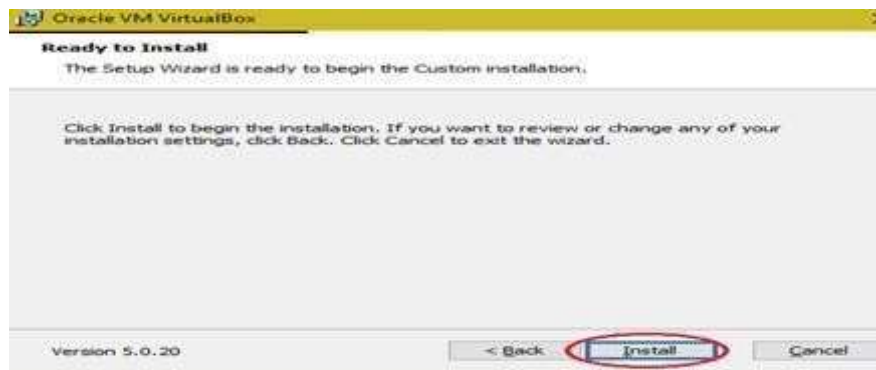
To write the procedure to install the virtualbox with different flavours of Linux/Windows.

**PROCEDURE:**

**Step 1:** Run the Virtualbox setup and click the “Next” option.



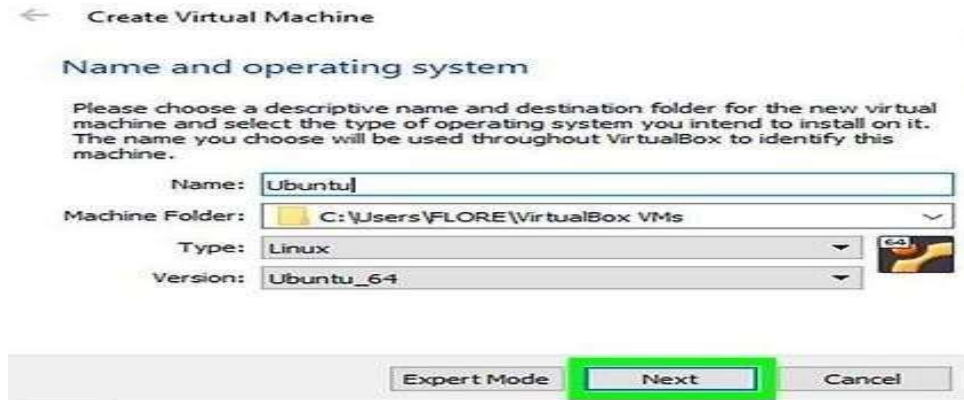
**Step 2:** Repeat Click on “Next” option and Finally Click the option “Install”



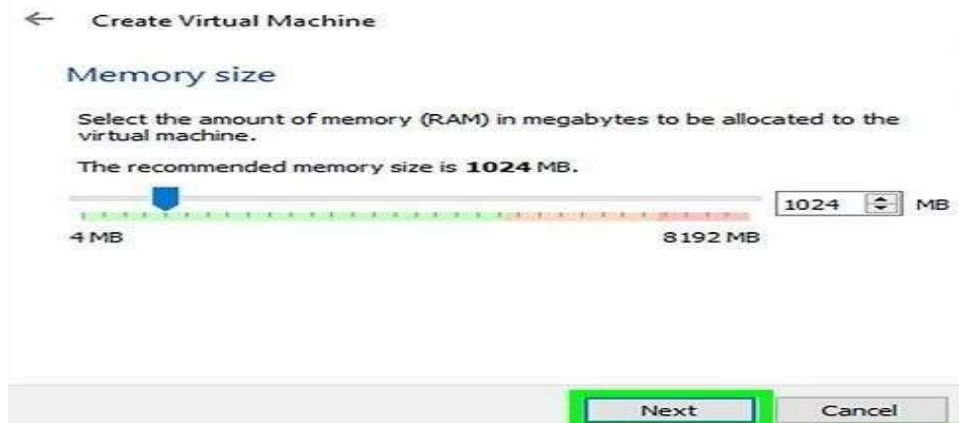
**Step 3:** To Install “Ubuntu” as virtual machine in “Oracle VM Virtualbox”, Open the Oracle VM Virtualbox Manager



#### Step 4: Create new virtual machine



#### Step 5: Select the memory size



#### Step 6: Create a virtual hard drive



## Step 7: Select the type of hardware file



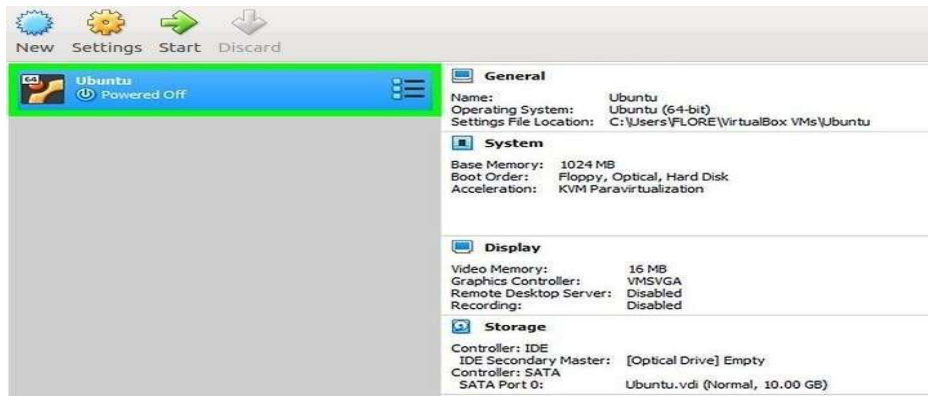
## Step 8: Select the type of storage on physical hard drive



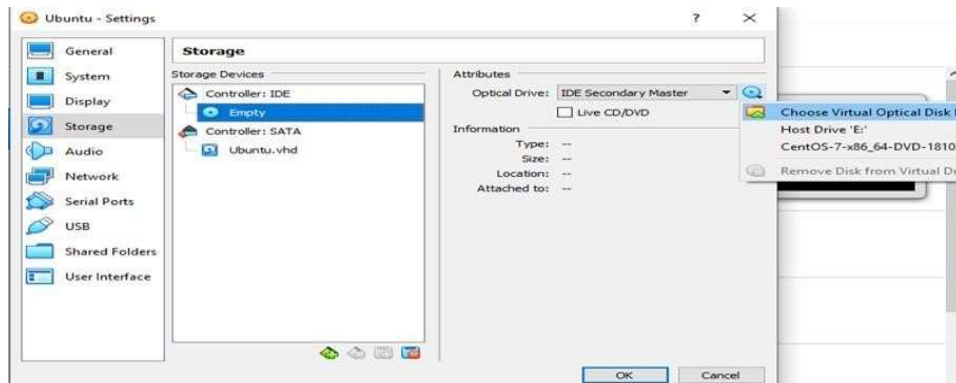
## Step 9: Select the size of virtual hard drive



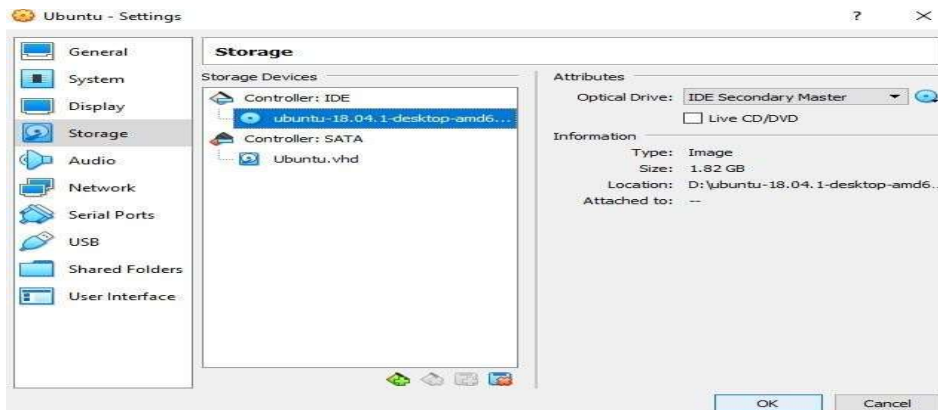
**Step 10:** Select the new virtual OS created and click on “Settings”



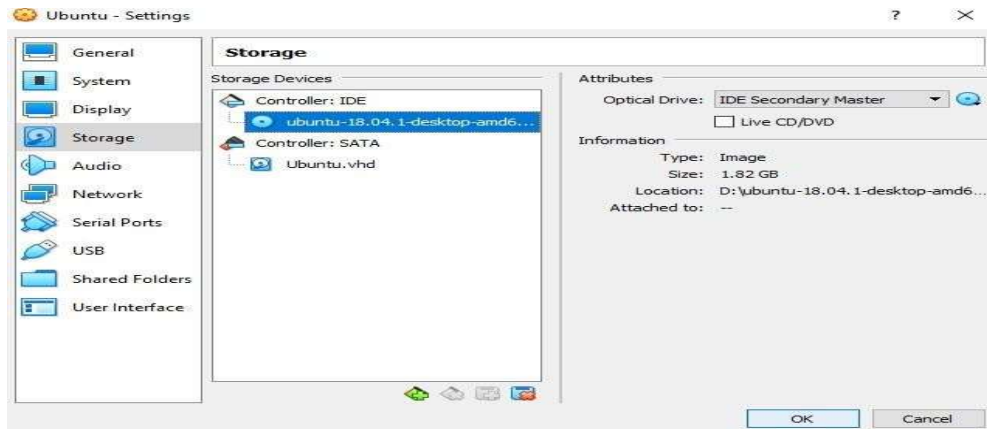
**Step 11:** Select “Storage” from the left panel of the window



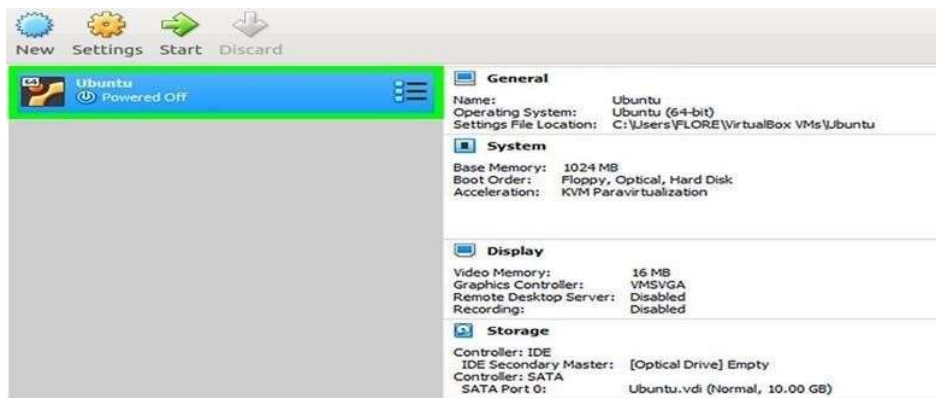
**Step 12:** Click on the first icon “Add CD/DVD device” in Controller:IDE



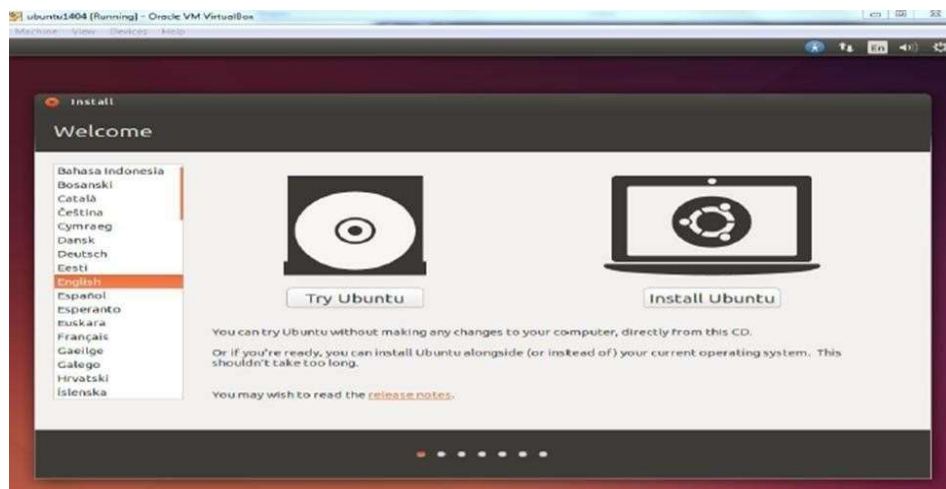
**Step 13:** Select “Choose Disk” and Choose the virtual machine to be used and click “Open”




**Step 14:** Click “OK” and select “Start” to run the virtual machine



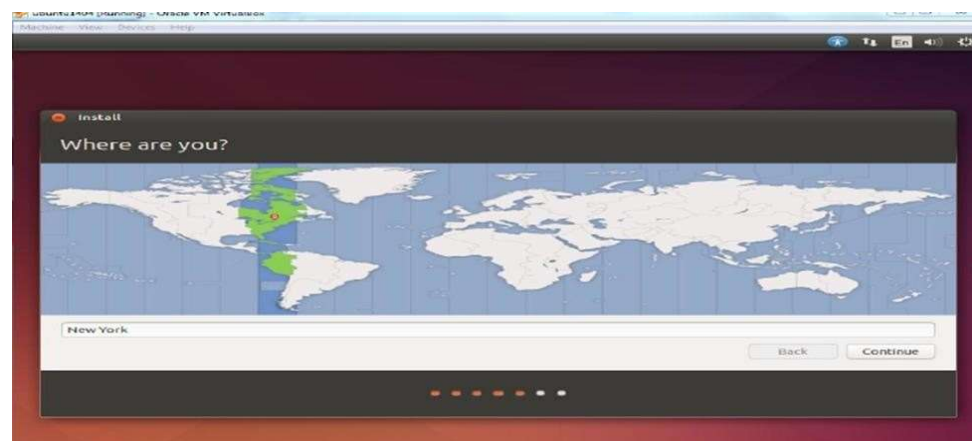
**Step 15:** To install ubuntu, Click 'Install Ubuntu' button



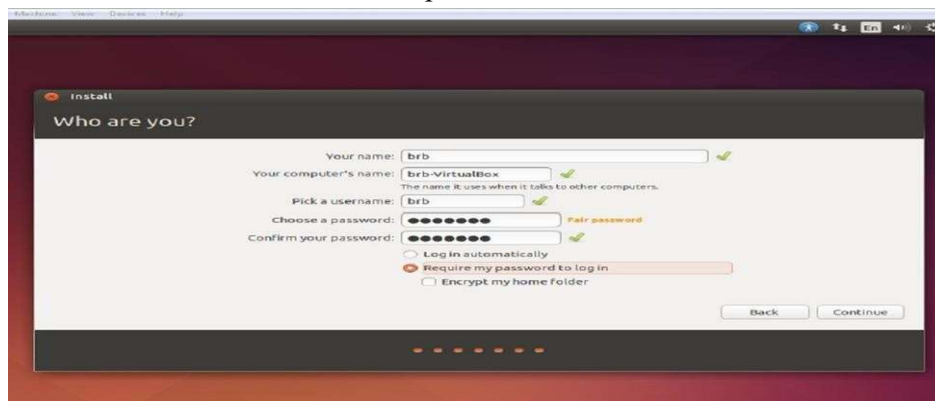
**Step 16:** Select 'Erase disk and install Ubuntu' option is selected and click 'Install Now' button



**Step 17:** Click 'Continue' button for upcoming dialogue box




**Step 18:** Enter the name, username and password. Click 'Continue' button






**Step 19:** After completion of installation process, click on 'Restart Now' button.



**Step 20:** The same way, we can install windows OS .

**OUTPUT:**

(i) Ubuntu Operating System in Virtual Machine



(ii) Windows7 Operating System in Virtual Machine



**RESULT:**

Thus the virtualbox with different flavours of Linux/Windows has been Installed Sucessfully.



**EX.NO: 2**

## **INSTALLATION OF A C COMPILER IN THE VIRTUAL MACHINE AND EXECUTING A SIMPLE PROGRAM**

**DATE:**

**AIM:**


To install a C compiler in the virtual machine and execute a simple program.

### **PROCEDURE:**

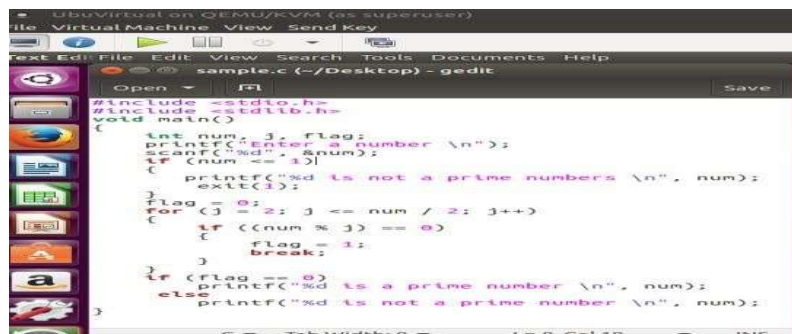
**STEP:1** Before install a C compiler in a virtual machine, we have Create a virtual machine by opening the Kernal Virtual Machine (KVM). Open the installed virtual manager. As well as install different Ubuntu and Windows OS in that virtual machine with different names.



**STEP:2** Open the Ubuntu OS in our Virtual Machine.



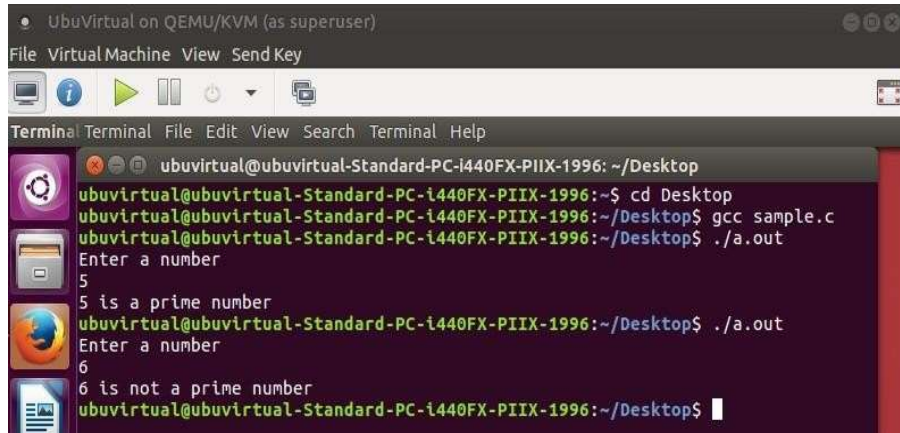
**STEP:3** Open TextEditor in Ubuntu OS and type a C program and save it in desktop to execute.



**STEP:4** To install the C compiler in Ubuntu OS, open the terminal and type the command.

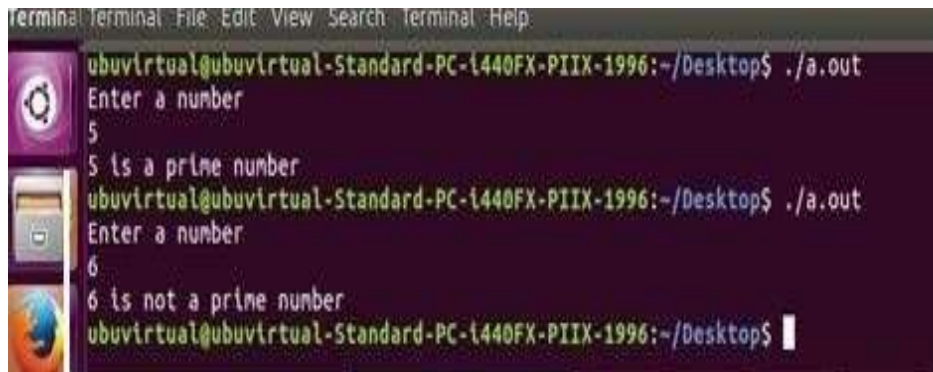
**\$sudo apt-get install gcc**

**STEP:5** And then compile the C program and execute it.



```
UbuVirtual on QEMU/KVM (as superuser)
File Virtual Machine View Send Key
Terminal Terminal File Edit View Search Terminal Help
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996: ~/Desktop
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~$ cd Desktop
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$ gcc sample.c
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$ ./a.out
Enter a number
5
5 is a prime number
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$ ./a.out
Enter a number
6
6 is not a prime number
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$
```

**OUTPUT:**



```
Terminal Terminal File Edit View Search Terminal Help
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$ ./a.out
Enter a number
5
5 is a prime number
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$ ./a.out
Enter a number
6
6 is not a prime number
ubuvirtual@ubuvirtual-Standard-PC-i440FX-PIIX-1996:~/Desktop$
```

**RESULT:**

Thus the C compiler is installed in the virtual machine and executed the program Successfully.

**EX.NO: 3**

## **INSTALLATION OF GOOGLE APP ENGINE AND CREATE HELLOWORLD APP USING JAVA**

**DATE:**

**AIM:**

To install google app engine and create hello world app using java.

**PROCEDURE:**

- **Create a project**

Projects bundle code, VMs, and other resources together for easier development and monitoring.

- **Build and run your "Hello, world!" app**

You will learn how to run your app using Cloud Shell, right in your browser. At the end, you'll deploy your app to the web using the App Engine Maven plugin.

GCP (Google Cloud Platform) organizes resources into projects, which collect all of the related resources for a single application in one place.

Begin by creating a new project or selecting an existing project for this tutorial.

**Step1:**

Select a project, or create a new one

**Step2:**

**Using Cloud Shell**

Cloud Shell is a built-in command-line tool for the console. You're going to use Cloud Shell to deploy your app.

**Open Cloud Shell**

Open Cloud Shell by clicking the

**Activate Cloud Shell** button in the navigation bar in the upper-right corner of the console  
**Clone the sample code**

Use Cloud Shell to clone and navigate to the "Hello World" code. The sample code is cloned from your project repository to the Cloud Shell.

Note: If the directory already exists, remove the previous files before cloning:

```
rm -rf appengine-try-java
```

```
git clone \
```

```
https://github.com/GoogleCloudPlatform/appengine-try-java
```

```
cd appengine-try-java
```

### **Step3:**

#### **Configuring your deployment**

You are now in the main directory for the sample code. You'll look at the files that configure your application.

#### **Exploring the application**

Enter the following command to view your application code:

```
cat \  
  
src/main/java/myapp/DemoServlet.java
```

This servlet responds to any request by sending a response containing the message Hello, world!.

#### **Exploring your configuration**

For Java, App Engine uses XML files to specify a deployment's configuration.

Enter the following command to view your configuration file:

```
cat pom.xml
```

The helloworld app uses Maven, which means you must specify a Project Object Model, or POM, which contains information about the project and configuration details used by Maven to build the project.

### **Step4:**

#### **Testing your app**

##### **Test your app on Cloud Shell**

Cloud Shell lets you test your app before deploying to make sure it's running as intended, just like debugging on your local machine.

To test your app enter the following:

```
mvn appengine:run
```

##### **Preview your app with "Web preview"**

Your app is now running on Cloud Shell. You can access the app by clicking the **Web preview**

button at the top of the Cloud Shell pane and choosing **Preview on port 8080**.

## Terminating the preview instance

Terminate the instance of the application by pressing Ctrl+C in the Cloud Shell

### Step5:

## Deploying to App Engine

### Create an application

To deploy your app, you need to create an app in a region:

```
gcloud app create
```

Note: If you already created an app, you can skip this step.

### Deploying with Cloud Shell

Now you can use Cloud Shell to deploy your app.

First, set which project to use:

```
gcloud config set project \  
<YOUR-PROJECT>
```

Then deploy your app:

```
mvn appengine:deploy
```

### Visit your app

Congratulations! Your app has been deployed.

The default URL of your app is a subdomain on appspot.com that starts with your project's ID: [.<your-project>.appspot.com](http://<your-project>.appspot.com).

Try visiting your deployed application.

### View your app's status

You can check in on your app by monitoring its status on the App Engine dashboard.

Open the **Navigation menu** in the upper-left corner of the console.

Then, select the **App Engine** section

### RESULT:

Thus google app engine is installed and hello world app using java created successfully.

## **EX.NO: 4    LAUNCH THE WEB APPLICATIONS USING GAE LAUNCHER**

**DATE:**

### **AIM:**

To launch the web applications using gae launcher.

### **PROCEDURE:**

Before you can host your website on Google App Engine:

1. Create a new Cloud Console project or retrieve the project ID of an existing project to use:

Go to the Projects page (<https://console.cloud.google.com/project>)

Tip: You can retrieve a list of your existing project IDs with the gcloud command line tool (#before\_you\_begin).

2. Install and then initialize the Google Cloud SDK:

Download the SDK (/sdk/docs)

Creating a website to host on Google App Engine

Basic structure for the project

This guide uses the following structure for the project:

app.yaml: Configure the settings of your App Engine application.

www/: Directory to store all of your static files, such as HTML, CSS, images, and JavaScript.

css/: Directory to store stylesheets.

style.css: Basic stylesheet that formats the look and feel of your site.

images/: Optional directory to store images.

index.html: An HTML file that displays content for your website.

js/: Optional directory to store JavaScript files.

Other asset directories.

Creating the app.yaml file

The app.yaml file is a configuration file that tells App Engine how to map URLs to your static files. In the following steps, you will add handlers that will load www/index.html when someone visits your website, and all static files will be stored in and called from the

www directory.

Create the app.yaml file in your application's root directory:

1. Create a directory that has the same name as your project ID. You can find your project ID in the Console (<https://console.cloud.google.com/>).
2. In directory that you just created, create a file named app.yaml.
3. Edit the app.yaml file and add the following code to the file:

```
runtime: python27
api_version: 1
threadsafe: true
handlers:
- url: /
  static_files: www/index.html
  upload: www/index.html
- url: /(.*)
  static_files: www/^1
  upload: www/(.*)
```

More reference information about the app.yaml file can be found in the app.yaml reference documentation (</appengine/docs/standard/python/config/appref>).

Creating the index.html \_le

Create an HTML file that will be served when someone navigates to the root page of your website. Store this file in your www directory.

Deploying your application to App Engine

When you deploy your application files, your website will be uploaded to App Engine. To deploy your app, run the following command from within the root directory of your application where the app.yaml file is located:

Optional flags:

Include the --project flag to specify an alternate Cloud Console project ID to what you initialized as the default in the gcloud tool. Example: --project [YOUR\_PROJECT\_ID]

Include the -v flag to specify a version ID, otherwise one is generated for you. Example: -v [YOUR\_VERSION\_ID]

|>

ead>

<title>Hello, world!</title>

<link rel="stylesheet" type="text/css" href="/css/style.css">



```
head>
ody>
<h1>Hello, world!</h1>
<p>
This is a simple static HTML file that will be served from Google App
Engine.
</p>
body>
ml>
```

### Deploying your application to App Engine

When you deploy your application files, your website will be uploaded to App Engine. To deploy your app, run the following command from within the root directory of your application where the app.yaml file is located:

```
id app deploy
```

Optional flags:

Include the --project flag to specify an alternate Cloud Console project ID to what you initialized as the default in the gcloud tool. Example: --project [YOUR\_PROJECT\_ID]

Include the -v flag to specify a version ID, otherwise one is generated for you. Example: -v [YOUR\_VERSION\_ID]

To learn more about deploying your app from the command line, see [Deploying a Python 2 App \(/appengine/docs/python/tools/uploadinganapp\)](#).

Viewing your application

To launch your browser and view the app at [https://PROJECT\\_ID.REGION\\_ID\(#appengine-urls\).r.appspot.com](https://PROJECT_ID.REGION_ID(#appengine-urls).r.appspot.com), run the following command

```
ud app browse
```

### RESULT:

Thus the web applications using gae launcher launched.

## **EX.NO: 5 SIMULATE A CLOUD SCENARIO USING CLOUDSIM AND RUN A SCHEDULING ALGORITHM THAT IS NOT PRESENT IN CLOUDSIM**

**DATE:**

### **What is Cloudsim?**

CloudSim is a simulation toolkit that supports the modelling and simulation of the core functionality of cloud, like job/task queue, processing of events, creation of cloud entities (datacentre, datacentre brokers, etc), communication between different entities, implementation of broker policies, etc. This toolkit allows to:

- Test application services in a repeatable and controllable environment.
- Tune the system bottlenecks before deploying apps in an actual cloud.
- Experiment with different workload mix and resource performance scenarios on simulated infrastructure for developing and testing adaptive application provisioning techniques

Core features of CloudSim are:

- The Support of modelling and simulation of large scale computing environment as federated cloud data centres, virtualized server hosts, with customizable policies for provisioning host resources to virtual machines and energy-aware computational resources
- It is a self-contained platform for modelling cloud's service brokers, provisioning, and allocation policies.
- It supports the simulation of network connections among simulated system elements.
- Support for simulation of federated cloud environment, that inter-networks resources from both private and public domains.
- Availability of a virtualization engine that aids in the creation and management of multiple independent and co-hosted virtual services on a data centre node.
- Flexibility to switch between space shared and time-shared allocation of processing cores to virtualized services.

```
import java.text.DecimalFormat;
```

```
import java.util.Calendar;
```

```
import java.util.List;
```

```
import org.cloudbus.cloudsim.Cloudlet;
```

```
import org.cloudbus.cloudsim.Datacenter;
```

```

import org.cloudbus.cloudsim.Log;
import org.cloudbus.cloudsim.Vm;
import org.cloudbus.cloudsim.core.CloudSim;

/**
 * FCFS Task scheduling
 * @author Linda J
 */
public class FCFS {

    /** The cloudlet list. */
    private static List<Cloudlet> cloudletList;

    /** The vmlist. */
    private static List<Vm> vmlist;

    private static int reqTasks = 5;
    private static int reqVms = 2;

    /**
     * Creates main() to run this example
     */
    public static void main(String[] args) {

        Log.println("Starting FCFS...");

        try {
            // First step: Initialize the CloudSim package. It should be called
            // before creating any entities.
            int num_user = 1; // number of cloud users
            Calendar calendar = Calendar.getInstance();
            boolean trace_flag = false; // mean trace events

            // Initialize the CloudSim library
            CloudSim.init(num_user, calendar, trace_flag);

```

```
// Second step: Create Datacenters  
//Datacenters are the resource providers in CloudSim. We need at list one of  
them to run a CloudSim simulation  
@SuppressWarnings("unused")  
Datacenter datacenter0 = createDatacenter("Datacenter_0");  
  
//Third step: Create Broker  
FcfsBroker broker = createBroker();  
int brokerId = broker.getId();  
  
//Fourth step: Create one virtual machine  
vmList = new VmsCreator().createRequiredVms(reqVms, brokerId);  
  
//submit vm list to the broker  
broker.submitVmList(vmList);  
  
//Fifth step: Create two Cloudlets  
cloudletList = new CloudletCreator().createUserCloudlet(reqTasks, brokerId);  
  
//submit cloudlet list to the broker  
broker.submitCloudletList(cloudletList);  
  
//call the scheduling function via the broker  
broker.scheduleTaskstoVms();  
  
// Sixth step: Starts the simulation  
CloudSim.startSimulation();  
  
// Final step: Print results when simulation is over  
List<Cloudlet> newList = broker.getCloudletReceivedList();  
  
CloudSim.stopSimulation();  
  
printCloudletList(newList);
```

```

        Log.println("FCFS finished!");
    }
    catch (Exception e) {
        e.printStackTrace();
        Log.println("The simulation has been terminated due to an unexpected error");
    }
}

private static Datacenter createDatacenter(String name){
    Datacenter datacenter=new DataCenterCreator().createUserDatacenter(name,
reqVms);

    return datacenter;

}

//We strongly encourage users to develop their own broker policies, to submit vms and
cloudlets according
//to the specific rules of the simulated scenario
private static FcfsBroker createBroker(){

    FcfsBroker broker = null;
    try {
        broker = new FcfsBroker("Broker");
    } catch (Exception e) {
        e.printStackTrace();
        return null;
    }
    return broker;
}

/**
 * Prints the Cloudlet objects
 * @param list list of Cloudlets
 */

```

```

private static void printCloudletList(List<Cloudlet> list) {
    int size = list.size();
    Cloudlet cloudlet;

    String indent = "  ";
    Log.println();
    Log.println("===== OUTPUT =====");
    Log.println("Cloudlet ID" + indent + "STATUS" + indent +
        "Data center ID" + indent + "VM ID" + indent + "Time" + indent + "Start Time"
+ indent + "Finish Time");

    DecimalFormat dft = new DecimalFormat("###.##");
    for (int i = 0; i < size; i++) {
        cloudlet = list.get(i);
        Log.print(indent + cloudlet.getCloudletId() + indent + indent);

        if (cloudlet.getCloudletStatus() == Cloudlet.SUCCESS){
            Log.print("SUCCESS");

            Log.println( indent + indent + cloudlet.getResourceId() + indent + indent +
indent + cloudlet.getVmId() +
                indent + indent + dft.format(cloudlet.getActualCPUTime()) + indent + indent
+ dft.format(cloudlet.getExecStartTime())+
                indent + indent + dft.format(cloudlet.getFinishTime()));
        }
    }
}

```

## RESULTS :

Thus simulating a cloud scenario using cloudsim is simulated successfully.

## EX.NO: 6      PROCEDURE TO TRANSFER THE FILES FROM ONE VIRTUAL MACHINE TO ANOTHER VIRTUAL MACHINE

**DATE:**

**Aim:**

To write a procedure to transfer the files from one Virtual Machine to another Virtual Machine.


**PROCEDURE:**

**STEP:1** Select the VM and click File->Export Appliance



**STEP:2** Select the VM to be exported and click NEXT.

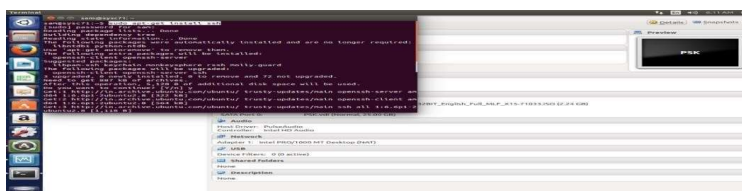
**STEP:3** Note the file path and click “Next”



**STEP:4** Click “Export”, The Virtual machine is being exported.



**STEP:5** Install “ssh” to access the neighbour's VM.






**STEP:6** Go to File->Computer:/home/sam/Documents/



**STEP:7** Type the neighbour's URL: sftp://sam@172.16.42.\_/




**STEP:8** Give the password and get connected.



**STEP:9** Select the VM and copy it in desktop.



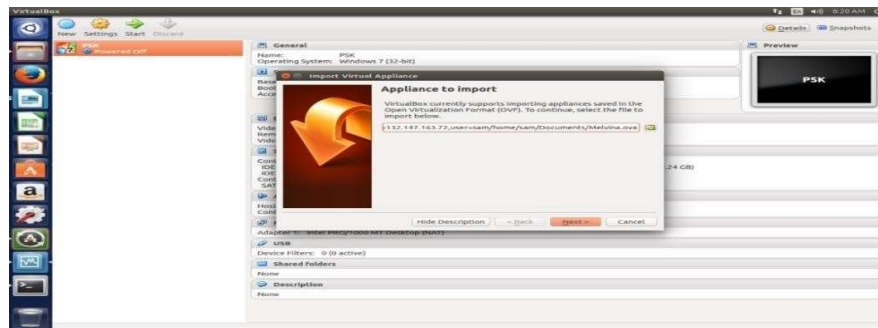
**STEP:10** Open VirtualBox and select File->Import Appliance->Browse




**STEP:11** Select the VM to be imported and click “Open”.



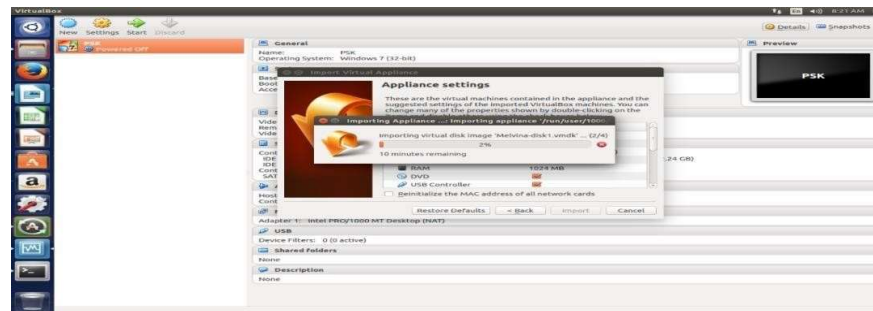
**STEP:12** Click “Next”



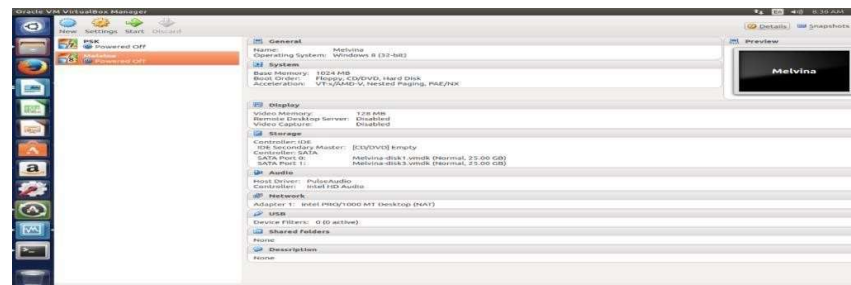
**STEP:13** Click “Import”.



**STEP:14** VM is being imported.



**STEP:15** VM is imported.



**RESULT:**

Thus the files from one Virtual Machine to another Virtual Machine is transferred Successfully..

## EXP NO :7 INSTALL HADOOP SINGLE NODE CLUSTER AND RUN SIMPLE APPLICATIONS LIKE WORDCOUNT

**DATE:**

**Aim:**

To Install Hadoop single node cluster and run simple applications like wordcount.

**Steps:**

### Install Hadoop

**Step 1:** [Click here](#) to download the Java 8 Package. Save this file in your home directory.

**Step 2:** Extract the Java Tar File.

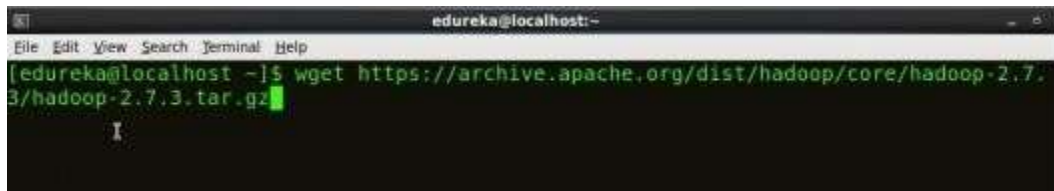
**Command:** `tar -xvf jdk-8u101-linux-i586.tar.gz`

A terminal window titled 'edureka@localhost:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The command '[edureka@localhost ~]\$ tar -xvf jdk-8u101-linux-i586.tar.gz' is entered and executed, with a green cursor at the end of the line.

*Fig: Hadoop Installation – Extracting Java Files*

**Step 3:** Download the Hadoop 2.7.3 Package.

**Command:** `wget https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz`

A terminal window titled 'edureka@localhost:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The command '[edureka@localhost ~]\$ wget https://archive.apache.org/dist/hadoop/core/hadoop-2.7.3/hadoop-2.7.3.tar.gz' is entered and executed, with a green cursor at the end of the line.

*Fig: Hadoop Installation – Downloading Hadoop*

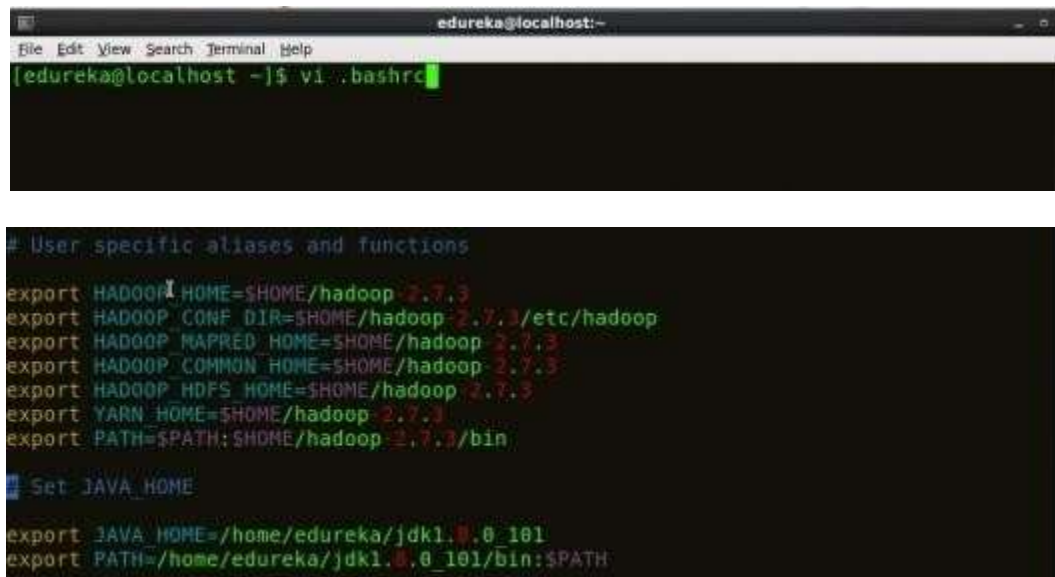
**Step 4:** Extract the Hadoop tar File.

**Command:** `tar -xvf hadoop-2.7.3.tar.gz`

A terminal window titled 'edureka@localhost:~ (on localhost.localdomain)' with a menu bar (File, Edit, View, Search, Terminal, Help). The command '[edureka@localhost ~]\$ tar -xvf hadoop-2.7.3.tar.gz' is entered and executed, with a green cursor at the end of the line.

**5:** Add the Hadoop and Java paths in the bash file (.bashrc). Open. **bashrc** file. Now, add Hadoop and Java Path as shown below.

**Command:** vi .bashrc



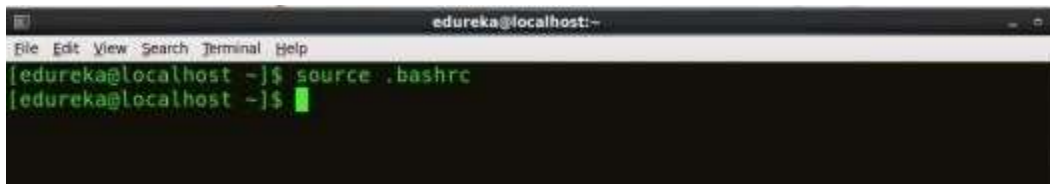
```
edureka@localhost:~  
File Edit View Search Terminal Help  
[edureka@localhost ~]$ vi .bashrc  
  
# User specific aliases and functions  
  
export HADOOP_HOME=$HOME/hadoop-2.7.3  
export HADOOP_CONF_DIR=$HOME/hadoop-2.7.3/etc/hadoop  
export HADOOP_MAPRED_HOME=$HOME/hadoop-2.7.3  
export HADOOP_COMMON_HOME=$HOME/hadoop-2.7.3  
export HADOOP_HDFS_HOME=$HOME/hadoop-2.7.3  
export YARN_HOME=$HOME/hadoop-2.7.3  
export PATH=$PATH:$HOME/hadoop-2.7.3/bin  
  
# Set JAVA_HOME  
  
export JAVA_HOME=/home/edureka/jdk1.8.0_101  
export PATH=/home/edureka/jdk1.8.0_101/bin:$PATH
```

*Fig: Hadoop Installation – Setting Environment Variable*

Then, save the bash file and close it.

For applying all these changes to the current Terminal, execute the source command.

**Command:** source .bashrc



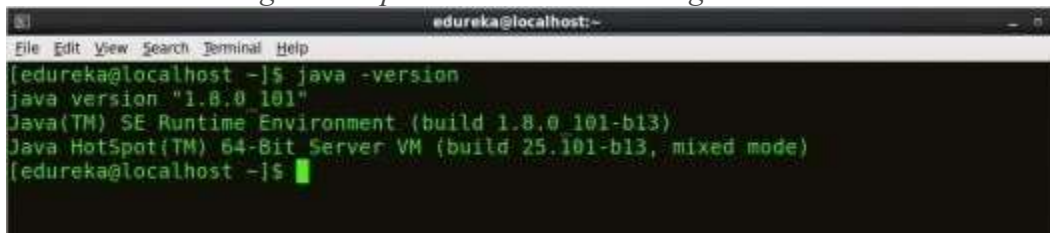
```
edureka@localhost:~  
File Edit View Search Terminal Help  
[edureka@localhost ~]$ source .bashrc  
[edureka@localhost ~]$
```

*Fig: Hadoop Installation – Refreshing environment variables*

To make sure that Java and Hadoop have been properly installed on your system and can be accessed through the Terminal, execute the java -version and hadoop version commands.

**Command:** java -version

*Fig: Hadoop Installation – Checking Java Version*



```
edureka@localhost:~  
File Edit View Search Terminal Help  
[edureka@localhost ~]$ java -version  
java version "1.8.0_101"  
Java(TM) SE Runtime Environment (build 1.8.0_101-b13)  
Java HotSpot(TM) 64-Bit Server VM (build 25.101-b13, mixed mode)  
[edureka@localhost ~]$
```

**Command:** `hadoop version`

```
edureka@localhost:~$ hadoop version
Hadoop 2.7.3
Subversion https://git-wip-us.apache.org/repos/asf/hadoop.git -r baa91f7c6bc9cb92be5982de4719c1c8af91ccff
Compiled by root on 2016-08-18T01:41Z
Compiled with protoc 2.5.0
From source with checksum 2e4ce5f957ea4db193bce3734ff29ff4
This command was run using /home/edureka/hadoop-2.7.3/share/hadoop/common/hadoop-common-2.7.3.jar
[edureka@localhost ~]$
```

*Fig: Hadoop Installation – Checking Hadoop Version*

**Step 6:** Edit the **Hadoop Configuration files**.

**Command:** `cd hadoop-2.7.3/etc/hadoop/`



**Command:** `ls`

All the Hadoop configuration files are located in **hadoop-2.7.3/etc/hadoop** directory as you can see in the snapshot below:

```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
[edureka@localhost ~]$ cd hadoop-2.7.3/etc/hadoop/
[edureka@localhost hadoop]$ ls
capacity-scheduler.xml      httpfs-env.sh              mapred-env.sh
configuration.xml          httpfs-log4j.properties   mapred-queues.xml.template
container-executor.cfg     httpfs-signature.secret   mapred-site.xml.template
core-site.xml              httpfs-site.xml           slaves
hadoop-env.cmd             kms-acls.xml              ssl-client.xml.example
hadoop-env.sh              kms-env.sh                ssl-server.xml.example
hadoop-metrics2.properties kms-log4j.properties     yarn-env.cmd
hadoop-metrics.properties kms-site.xml              yarn-env.sh
hadoop-policy.xml          log4j.properties         yarn-site.xml
hdfs-site.xml              mapred-env.cmd
```

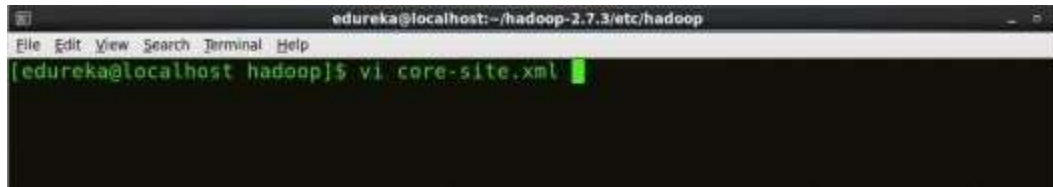
*Fig: Hadoop Installation – Hadoop Configuration Files*



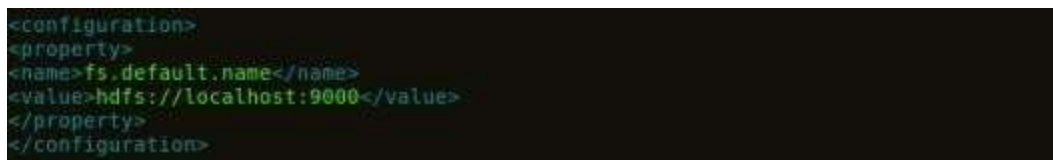
**Step 7:** Open *core-site.xml* and edit the property mentioned below inside configuration tag:

*core-site.xml* informs Hadoop daemon where NameNode runs in the cluster. It contains configuration settings of Hadoop core such as I/O settings that are common to HDFS & MapReduce.

**Command:** vi core-site.xml

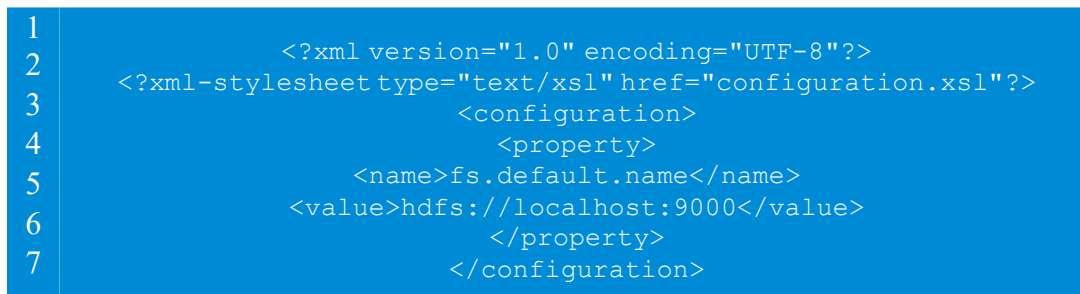


```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi core-site.xml
```



```
<configuration>
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
</configuration>
```

*Fig: Hadoop Installation – Configuring core-site.xml*

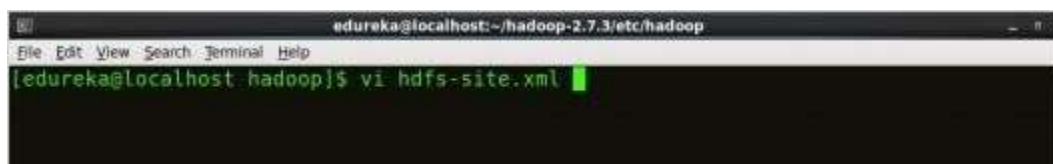


```
1      <?xml version="1.0" encoding="UTF-8"?>
2      <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3          <configuration>
4              <property>
5                  <name>fs.default.name</name>
6                  <value>hdfs://localhost:9000</value>
7              </property>
            </configuration>
```

**Step 8:** Edit *hdfs-site.xml* and edit the property mentioned below inside configuration tag:

*hdfs-site.xml* contains configuration settings of HDFS daemons (i.e. NameNode, DataNode, Secondary NameNode). It also includes the replication factor and block size of HDFS.

**Command:** vi hdfs-site.xml



```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi hdfs-site.xml
```



```
<configuration>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.permission</name>
<value>>false</value>
</property>
```



*Fig: Hadoop Installation – Configuring `hdfs-site.xml`*

```
1
2      <?xml version="1.0" encoding="UTF-8"?>
3  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
4      <configuration>
5          <property>
6              <name>dfs.replication</name>
7              <value>1</value>
8          </property>
9          <property>
10             <name>dfs.permission</name>
11             <value>>false</value>
12         </property>
13     </configuration>
```

**Step 9:** Edit the *mapred-site.xml* file and edit the property mentioned below

inside configuration tag:

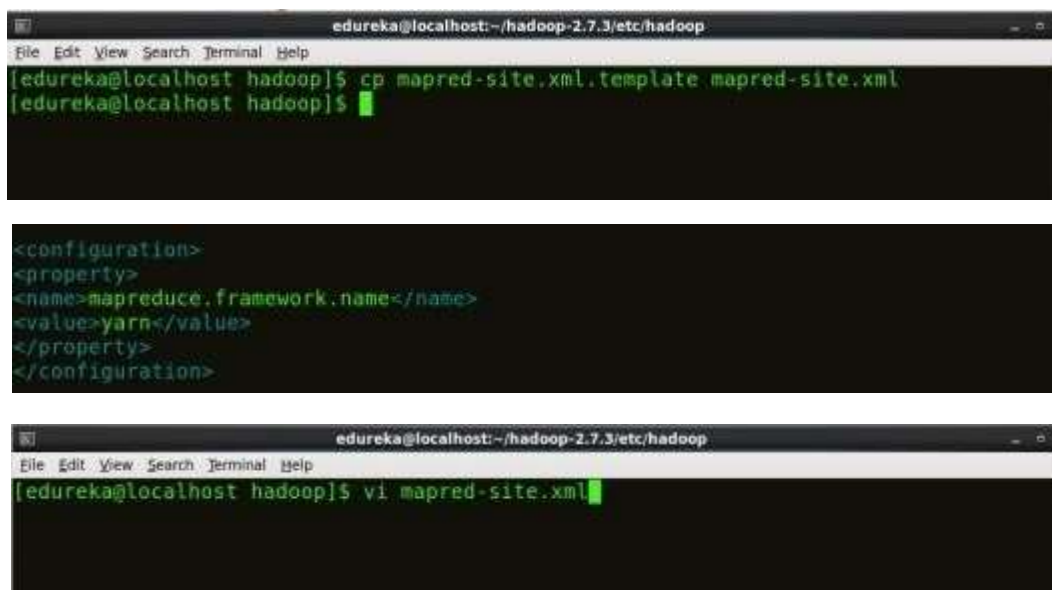
*mapred-site.xml* contains configuration settings of MapReduce application like number of JVM that can run in parallel, the size of the mapper and the reducer process, CPU cores available for a process, etc.

In some cases, *mapred-site.xml* file is not available. So, we have to create the *mapred-site.xml* file using *mapred-site.xml* template.

**Command:** `cp mapred-site.xml.template mapred-site.xml`

**Command:** `vi mapred-site.xml`.

*Fig: Hadoop Installation – Configuring `mapred-site.xml`*



The figure consists of three screenshots from a terminal window. The first screenshot shows the command `cp mapred-site.xml.template mapred-site.xml` being executed. The second screenshot shows the content of the newly created `mapred-site.xml` file, which contains a configuration tag with a property for `mapreduce.framework.name` set to `yarn`. The third screenshot shows the command `vi mapred-site.xml` being entered to edit the file.

```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ cp mapred-site.xml.template mapred-site.xml
[edureka@localhost hadoop]$

<configuration>
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
</configuration>

edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi mapred-site.xml
```

```

1      <?xml version="1.0" encoding="UTF-8"?>
2      <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3          <configuration>
4              <property>
5                  <name>mapreduce.framework.name</name>
6                  <value>yarn</value>
7              </property>
            </configuration>

```

**Step 10:** Edit *yarn-site.xml* and edit the property mentioned below inside configuration tag:

*yarn-site.xml* contains configuration settings of ResourceManager and NodeManager like application memory management size, the operation needed on program & algorithm, etc.

**Command:** vi yarn-site.xml



```

edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi yarn-site.xml

<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>

```

Fig: Hadoop Installation – Configuring yarn-site.xml

**Step 11:** Edit *hadoop-env.sh* and add the Java Path as mentioned below:

```

1
2      <?xml version="1.0">
3      <configuration>
4          <property>
5              <name>yarn.nodemanager.aux-services</name>
6              <value>mapreduce_shuffle</value>
7          </property>
8          <property>
9              <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</
10             name>
11             <value>org.apache.hadoop.mapred.ShuffleHandler</value>
            </property>

```

*hadoop-env.sh* contains the environment variables that are used in the script to run Hadoop like Java home path, etc.

**Command:** vi hadoop-env.sh

A terminal window titled 'edureka@localhost:~/hadoop-2.7.3/etc/hadoop' showing the command 'vi hadoop-env.sh' being executed. The terminal then shows the command 'export JAVA\_HOME=/home/edureka/jdk1.8.0\_101' being entered.

```
edureka@localhost:~/hadoop-2.7.3/etc/hadoop
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ vi hadoop-env.sh

# The java implementation to use.
export JAVA_HOME=/home/edureka/jdk1.8.0_101
```

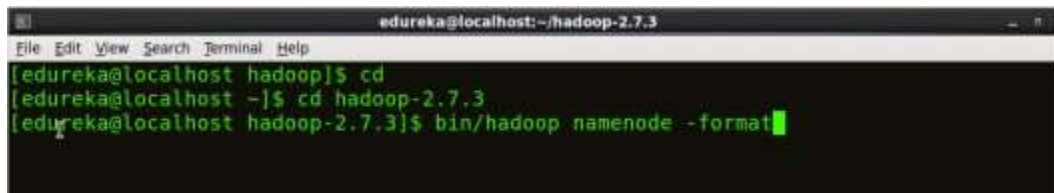
*Fig: Hadoop Installation – Configuring hadoop-env.sh Step*

**12: Go to Hadoop home directory and format the NameNode.**

**Command:** cd

**Command:** cd hadoop-2.7.3

**Command:** bin/hadoop namenode -format

A terminal window titled 'edureka@localhost:~/hadoop-2.7.3' showing the sequence of commands to format the NameNode: 'cd', 'cd hadoop-2.7.3', and 'bin/hadoop namenode -format'.

```
edureka@localhost:~/hadoop-2.7.3
File Edit View Search Terminal Help
[edureka@localhost hadoop]$ cd
[edureka@localhost ~]$ cd hadoop-2.7.3
[edureka@localhost hadoop-2.7.3]$ bin/hadoop namenode -format
```

*Fig: Hadoop Installation – Formatting NameNode*

This formats the HDFS via NameNode. This command is only executed for the first time. Formatting the file system means initializing the directory specified by the dfs.name.dir variable.

Never format, up and running Hadoop filesystem. You will lose all your data stored in the HDFS.

**Step 13: Once the NameNode is formatted, go to hadoop-2.7.3/sbin directory and start all the daemons.**

**Command:** cd hadoop-2.7.3/sbin

Either you can start all daemons with a single command or do it individually.

**Command:** ./start-all.sh

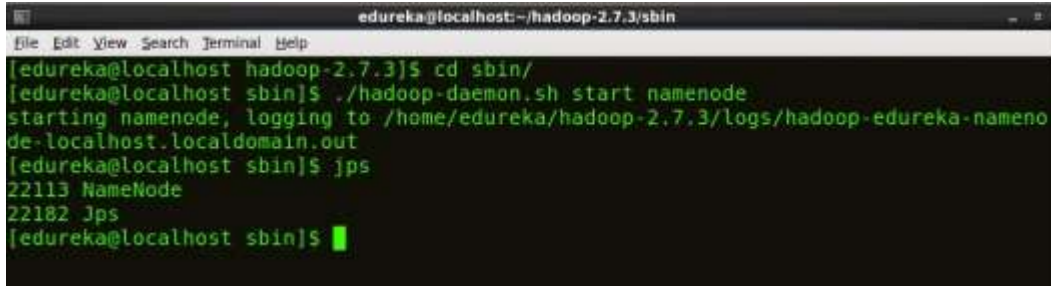
The above command is a combination of *start-dfs.sh*, *start-yarn.sh* & *mr-jobhistory-daemon.sh*

Or you can run all the services individually as below:

## Start NameNode:

The NameNode is the centerpiece of an HDFS file system. It keeps the directory tree of all files stored in the HDFS and tracks all the file stored across the cluster.

**Command:** `./hadoop-daemon.sh start namenode`

A terminal window titled 'edureka@localhost:~/hadoop-2.7.3/sbin' showing the execution of the command to start the NameNode. The user navigates to the 'sbin' directory and runs './hadoop-daemon.sh start namenode', which outputs 'starting namenode, logging to /home/edureka/hadoop-2.7.3/logs/hadoop-edureka-namenode-localhost.localdomain.out'. Then, the user runs 'jps', which outputs '22113 NameNode' and '22182 Jps'.

```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost hadoop-2.7.3]$ cd sbin/
[edureka@localhost sbin]$ ./hadoop-daemon.sh start namenode
starting namenode, logging to /home/edureka/hadoop-2.7.3/logs/hadoop-edureka-namenode-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22182 Jps
[edureka@localhost sbin]$
```

### Start DataNode:

On startup, a DataNode connects to the Namenode and it responds to the requests from the Namenode for different operations.

**Command:** `./hadoop-daemon.sh start datanode`



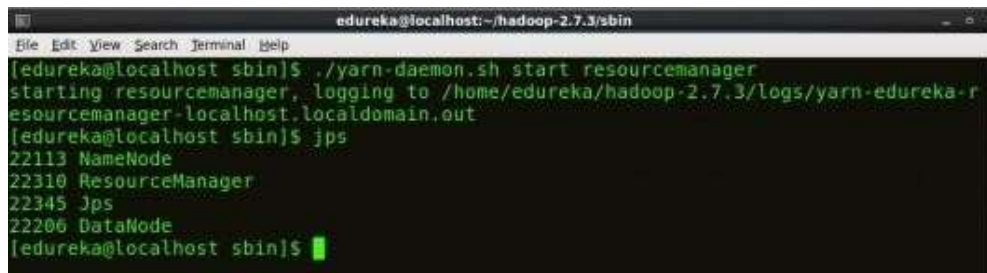
```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./hadoop-daemon.sh start datanode
starting datanode, logging to /home/edureka/hadoop-2.7.3/logs/hadoop-edureka-datano
de-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22278 Jps
22206 DataNode
[edureka@localhost sbin]$
```

*Fig: Hadoop Installation – Starting DataNode*

### Start ResourceManager:

ResourceManager is the master that arbitrates all the available cluster resources and thus helps in managing the distributed applications running on the YARN system. Its work is to manage each NodeManagers and the each application's ApplicationMaster.

**Command:** `./yarn-daemon.sh start resourcemanager`



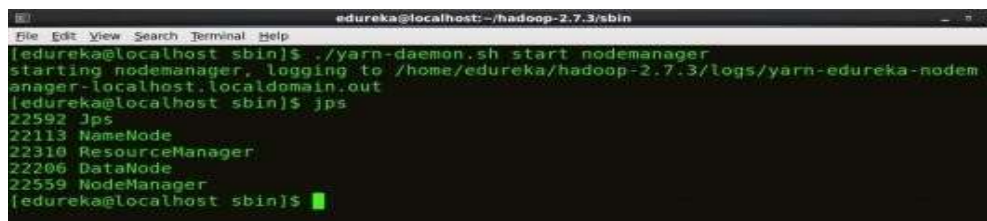
```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./yarn-daemon.sh start resourcemanager
starting resourcemanager, logging to /home/edureka/hadoop-2.7.3/logs/yarn-edureka-r
esourcemanager-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22310 ResourceManager
22345 Jps
22206 DataNode
[edureka@localhost sbin]$
```

*Fig: Hadoop Installation – Starting ResourceManager*

### Start NodeManager:

The NodeManager in each machine framework is the agent which is responsible for managing containers, monitoring their resource usage and reporting the same to the ResourceManager.

**Command:** `./yarn-daemon.sh start nodemanager`



```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./yarn-daemon.sh start nodemanager
starting nodemanager, logging to /home/edureka/hadoop-2.7.3/logs/yarn-edureka-nodem
anager-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22592 Jps
22113 NameNode
22310 ResourceManager
22206 DataNode
22559 NodeManager
[edureka@localhost sbin]$
```



[See Batch Details](#)

*Fig: Hadoop Installation – Starting NodeManager*

#### **Start JobHistoryServer:**

JobHistoryServer is responsible for servicing all job history related requests from client.

**Command:** `./mr-jobhistory-daemon.sh start historyserver`

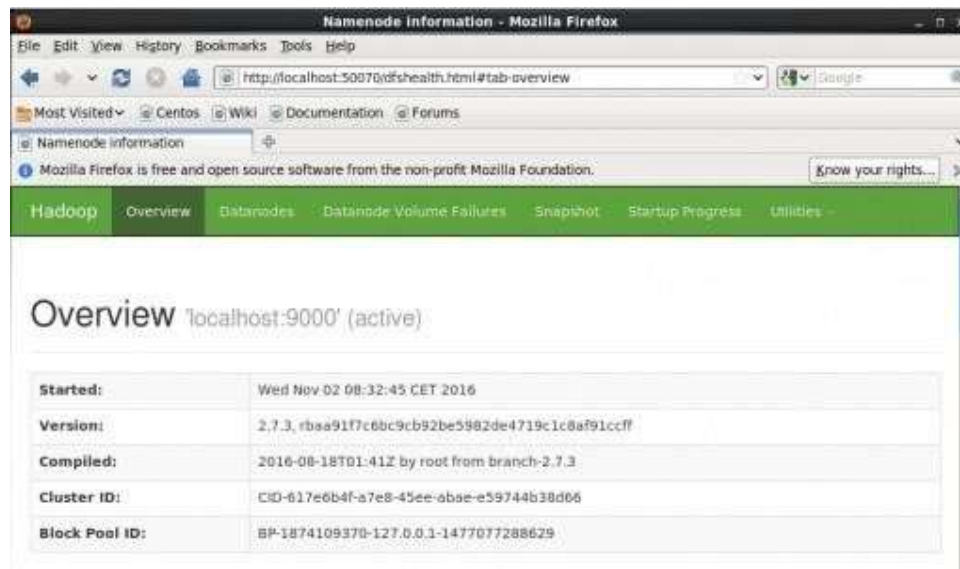
Step 14: To check that all the Hadoop services are up and running, run the below command.

**Command:** `jps`

```
edureka@localhost:~/hadoop-2.7.3/sbin
File Edit View Search Terminal Help
[edureka@localhost sbin]$ ./mr-jobhistory-daemon.sh start historyserver
starting historyserver, logging to /home/edureka/hadoop-2.7.3/logs/mapred-edureka-h
istoryserver-localhost.localdomain.out
[edureka@localhost sbin]$ jps
22113 NameNode
22310 ResourceManager
22694 JobHistoryServer
22727 Jps
22206 DataNode
22559 NodeManager
[edureka@localhost sbin]$
```

*Fig: Hadoop Installation – Checking Daemons*

**Step 15:** Now open the Mozilla browser and go to **localhost:50070/dfshealth.html** to check the NameNode interface.



*Fig: Hadoop Installation – Starting WebUI*

Congratulations, you have successfully installed a single node Hadoop cluster.

**Result:**

Thus the Hadoop one cluster was installed and simple applications executed successfully.



## **EXP NO :8 CREATING AND EXECUTING YOUR FIRST CONTAINER USING DOCKER.**

**DATE:**

**AIM:**

To create and execute your first container using Docker.

### **PROCEDURE**

To follow the instructions specific to your operating system. Here are the general steps for the most popular operating systems:

Windows:

- a. Visit the Docker website (<https://www.docker.com/>) and navigate to the "Get Docker" section.
- b. Click on the "Download for Windows" button to download the Docker Desktop installer.
- c. Run the installer and follow the prompts. During the installation, Docker may require you to enable Hyper-V and Containers features, so make sure to enable them if prompted.
- d. Once the installation is complete, Docker Desktop will be installed on your Windows machine. You can access it from the Start menu or the system tray.

Mac:

- a. Visit the Docker website (<https://www.docker.com/>) and navigate to the "Get Docker" section.
- b. Click on the "Download for Mac" button to download the Docker Desktop installer.
- c. Run the installer and drag the Docker icon to the Applications folder to install Docker Desktop.
- d. Launch Docker Desktop from the Applications folder or the Launchpad. It will appear in the status bar at the top of your screen.

Linux:

Docker supports various Linux distributions. The exact installation steps may vary based on your distribution. Here's a general outline:

- a. Visit the Docker website (<https://www.docker.com/>) and navigate to the "Get Docker" section.
- b. Click on the "Download for Linux" button.
- c. Docker provides installation instructions for various Linux distributions such as Ubuntu, CentOS, Debian, Fedora, and more. Follow the instructions specific to your distribution.
- d. Once Docker is installed, start the Docker service using the appropriate command for your Linux distribution.

After completing the installation, you can open a terminal or command prompt and run the `docker --version` command to verify that Docker is installed correctly. It should display the version of Docker installed on your system.

That's it! You now have Docker installed on your machine and can start using it to manage containers.

**Install Docker:** First, you need to install Docker on your machine. Docker provides platform-specific installation instructions on their website for different operating systems. Follow the instructions to install Docker for your particular OS.

**Docker Image:** Docker containers are created based on Docker images. An image is a lightweight, standalone, and executable package that includes everything needed to run a piece of software, including the code, runtime, libraries, and system tools. Docker Hub ([hub.docker.com](https://hub.docker.com)) is a popular online repository of Docker images. You can search for existing images on Docker Hub or create your own. For this example, we'll use an existing image.

**Pull an Image:** Open your terminal or command prompt and execute the following command to pull an existing Docker image from Docker Hub. We'll use the official hello-world image as an example:

Copy code

```
docker pull hello-world
```

Docker will download the image from the Docker Hub repository.

**Run a Container:** Once you have the Docker image, you can create and run a container based on that image. Execute the following command to run the hello-world container:

arduino

Copy code

```
docker run hello-world
```

Docker will create a container from the image and execute it. The container will print a "Hello from Docker!" message along with some information about your Docker installation.

**Note:** If you haven't pulled the hello-world image in the previous step, Docker will automatically download it before running the container.

**Congratulations!** You've created and executed your first Docker container. Docker will handle the container lifecycle, including starting, stopping, and managing resources for you. This simple example demonstrates the basic concept of running a container using Docker.

You can explore further by trying out different Docker images and running more complex applications within containers.

## **RESULT:**

Thus the first container using docker is created and executed successfully.

## EXP NO :9

## RUN A CONTAINER FROM DOCKER HUB

### AIM:

To run a container from docker hub

### PROCEDURE:

#### Run a Container from Docker Hub

To run a container from Docker Hub, you need to follow these steps:

**Search for an Image:** Visit the Docker Hub website (<https://hub.docker.com/>) and use the search bar to find the image you want to run. You can search for popular images like nginx, mysql, redis, etc., or specific images based on your requirements.

**Pull the Image:** Once you've found the desired image, open a terminal or command prompt and execute the following command to pull the image from Docker Hub:

php

Copy code

```
docker pull <image_name>
```

Replace <image\_name> with the name of the image you want to pull. For example, if you want to pull the nginx image, you would use:

Copy code

```
docker pull nginx
```

Docker will download the image and store it on your local machine.

**Run the Container:** After pulling the image, you can create and run a container based on that image. Use the following command:

arduino

Copy code

```
docker run <image_name>
```

Replace <image\_name> with the name of the image you pulled. For example:

arduino

Copy code

```
docker run nginx
```

Docker will create a container from the image and start it. The container will run the default command specified in the image, such as starting a web server, database, or any other application.

**Note:** By default, Docker will allocate a random port on your host machine and map it to the container's exposed ports. If you want to specify a specific port mapping, you can use the -p option. For example:

arduino

Copy code

```
docker run -p 8080:80 nginx
```

This will map port 8080 on your host machine to port 80 inside the container.

**Interact with the Container:** Once the container is running, you can interact with it as needed. For example, if you ran an nginx container, you can access it in your web browser by visiting <http://localhost> or [http://<your\\_host\\_ip>](http://<your_host_ip>) (if you specified a port mapping).

To stop the container, you can use the docker stop command followed by the container ID or name:

arduino

Copy code

docker stop <container\_id\_or\_name>

You can list the running containers using the docker ps command and stop them as necessary.

(or)

Flow-1: Pull Docker Image from Docker Hub and Run it

Step-1: Verify Docker version and also login to Docker Hub

docker version

docker login

Step-2: Pull Image from Docker Hub

docker pull stacksimplify/dockerintro-springboot-helloworld-rest-api:1.0.0-RELEASE

Step-3: Run the downloaded Docker Image & Access the Application

Copy the docker image name from Docker Hub

docker run --name app1 -p 80:8080 -d stacksimplify/dockerintro-springboot-helloworld-rest-api:1.0.0-RELEASE

http://localhost/hello

Step-4: List Running Containers

docker ps

docker ps -a

docker ps -a -q

Step-5: Connect to Container Terminal

docker exec -it <container-name> /bin/sh

Step-6: Container Stop, Start

docker stop <container-name>

docker start <container-name>

## RESULT:

Thus the program ran a container from docker hub

